

Bayesian Inference and Deep Learning

A/Prof Richard Yi Da Xu

Yida.Xu@uts.edu.au

Wechat: aubedata

<https://github.com/roboticcam/machine-learning-notes>

University of Technology Sydney (UTS)

March 3, 2021

preambles

In a Deep Learning dominated world ...

- ▶ Machine Learning has been dominated by Deep Learning since early 2010s
- ▶ many “hot” buzz-word Deep Learning frameworks has been over-studied:
 - ▶ various Neural Network architectures... Convolution, Recurrent, Residual, Self Attention, ...
 - ▶ Generative Adversarial networks, Normalizing Flow, Noise Contrast Estimation
 - ▶ Reinforcement Learning, Meta Learning, Transfer Learning
 - ▶ ...
- ▶ However, Deep Learning also has its drawbacks, they in general
 - ▶ require very large amount of data to perform better than other techniques
 - ▶ predominantly used for supervised learning, where correctness of labels are paramount
 - ▶ difficult to study the confidence of the result

$$\underbrace{p(\theta|X)}_{\text{posterior}} = \frac{\underbrace{p(X|\theta)}_{\text{likelihood}} \underbrace{p(\theta)}_{\text{prior}}}{\underbrace{p(X)}_{\text{normalization constant}}}$$
$$= \frac{p(X|\theta)p(\theta)}{\int_{\theta} p(X|\theta)p(\theta)}$$

- ▶ **interpretability** can determine a % credible interval of true parameter
- ▶ **natural way to incorporate prior information** for example, in sequential modeling, posterior of state at $t - 1$ becomes prior at time t
- ▶ **naturally applicable** to many unsupervised learning tasks, such as hierarchical models
- ▶ **readily available tools** many numerical methods such as MCMC available for inference
- ▶ ...

so the big question...

Can we bring the best of both worlds together?

1. Bayesian Inference **Tools**:

1.1 Variational inference

1.2 Monte-Carlo inference and their role in Deep Learning

1.2.1 **non-MCMC**:

Rejection, adaptive rejection, Importance, Sequential Monte Carlo,

1.2.2 **MCMC**:

Metropolis-Hasting, Slice, Swendsen-Wang, Hamiltonian Monte-Carlo

2. Where Bayesian Inference meets Deep Learning

2.1 Bayesian framework to **explain** Deep Learning

2.1.1 infinite-width Neural Network as Gaussian Process

2.2 Bayesian framework to **assist** Deep Learning

2.2.1 Bayesian assisted Generative Adversarial Networks

2.2.2 Determinantal Point Process assisted neural networks

2.2.3 Re-parameterization techniques

2.3 Bayesian framework to **be assist by** Deep Learning

2.3.1 Neural Networks assisted Recursive Bayesian Filtering

3. Final remark: where future is heading

► in this talk, I show a mixture of:

1. work by other researchers
2. our published and current work in this topic
3. making references to my **detailed tutorials**
4. no nice pictures, as I will hand-draw them using stylus

Part 1: Bayesian Inference Tools

1. Variational inference: one page
2. Monte-Carlo inference: discussed in detail
3. Laplace approximation: not discussed at all

Bayesian Inference Tools: Variational Inference

$$\begin{aligned}\ln(p(\mathbf{x})) &= \log \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z} \\ &= \log \int \frac{q(\mathbf{z}|\mathbf{x})}{q(\mathbf{z}|\mathbf{x})} p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z} \quad \text{can also be just } q(\mathbf{z}) \text{ approximate distribution} \\ &= \log \int \left(\frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} p(\mathbf{x}|\mathbf{z}) \right) q(\mathbf{z}|\mathbf{x})d\mathbf{z} \quad \text{re-arrange} \\ &\geq \int \log \left(\frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} p(\mathbf{x}|\mathbf{z}) \right) q(\mathbf{z}|\mathbf{x})d\mathbf{z} \quad \text{Jensen's equality} \\ &= \mathcal{L}(q) \quad \text{Evidence Lower BOund (ELBO)}\end{aligned}$$

► ELBO split **one**

$$\begin{aligned}&= \int \log p(\mathbf{x}|\mathbf{z})q(\mathbf{z}|\mathbf{x})d\mathbf{z} + \int \log \left(\frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \right) q(\mathbf{z}|\mathbf{x})d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}|\mathbf{z})] - \int \log \left(\frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z})} \right) q(\mathbf{z}|\mathbf{x})d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}|\mathbf{z})] - \text{KL}[q(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z})] \\ &= \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - \text{KL}[q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z})]\end{aligned}$$

► ELBO split **two**

$$\begin{aligned}&= \int \log p(\mathbf{x}|\mathbf{z})p(\mathbf{z})q(\mathbf{z}|\mathbf{x})d\mathbf{z} + \int \log \left(\frac{1}{q(\mathbf{z}|\mathbf{x})} \right) q(\mathbf{z}|\mathbf{x})d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}, \mathbf{z})] - \int \log q(\mathbf{z}|\mathbf{x})q(\mathbf{z}|\mathbf{x})d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}, \mathbf{z}) - \log q(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})]\end{aligned}$$

my tutorial "Variational Inference" <https://github.com/roboticcam/machine-learning-notes/blob/master/files/variational.pdf>

1. Non-MCMC:

Rejection, adaptive rejection, importance, Sequential Monte Carlo

my tutorial “Monte-Carlo method: an introduction”

https://github.com/roboticcam/machine-learning-notes/blob/master/files/introduction_monte_carlo.pdf

2. MCMC:

Metropolis-Hasting, Gibbs, Slice, Swendsen-Wang, Hamiltonian Monte-Carlo

my tutorial “Stochastic Matrices”

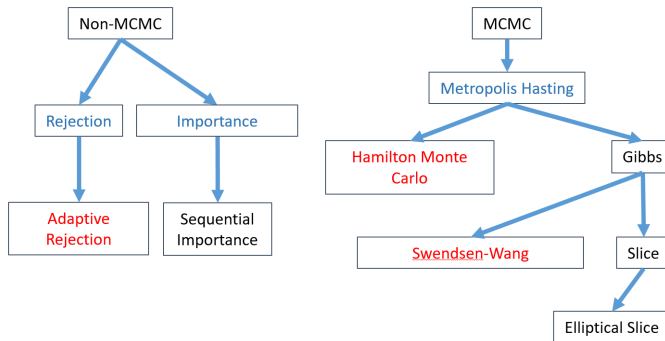
https://github.com/roboticcam/machine-learning-notes/blob/master/files/stochastic_matrices.pdf

my tutorial “Markov Chain Monte Carlo”


https://github.com/roboticcam/machine-learning-notes/blob/master/files/markov_chain_monte_carlo.pdf

Hierarchies of Monte Carlo Inference methods

there were a **subset** of methods either I used, or I wrote into my tutorial:



let me draw the ones in blue and show Python demo to ones in red



- ▶ Sampling is all about efficiency
- ▶ Rejection sampling can give you quite low acceptance ratio, should you choose a non-compatible $q(\cdot)$
- ▶ let's take a look at **adaptive rejection sampling**!

The algorithm

```
i = 0
while i ≠ N
  x(i) ~ q(x) and u ~ U(0, 1)
  if u <  $\frac{p(x(i))}{Mq(x(i))}$  then
    accept x(i)
    i = i + 1
  else
    reject x(i)
  end
end
```

Simplest MCMC algorithm: Metropolis Hasting

1. initialize $x^{(0)}$
2. **for** $i = 0$ to $N - 1$
 - $u \sim \mathbf{U}(0, 1)$
 - $x^* \sim q(x^* | x^{(i)})$
 - if** $u < \min \left(1, \frac{\pi(x^*)q(x | x^*)}{\pi(x)q(x^* | x)} \right)$
 - $x^{(i+1)} = x^*$
 - else**
 - $x^{(i+1)} = x^{(i)}$

- ▶ The take-home message here, is that it does not “discard” samples like rejection sampling. It simply “repeats” samples.
- ▶ If the same sample repeats too many times, it has **bad mixing**
- ▶ the key: to make M-H more efficient algorithm, we need to:

$$\min \left(1, \frac{\pi(x^*)q(x | x^*)}{\pi(x)q(x^* | x)} \right) \rightarrow 1$$

Hamiltonian Metropolis Hasting (HMC): very brief introduction

- ▶ Hamiltonian = Potential + Kinetic

$$H(q, p) = U(q) + K(p)$$

- ▶ let q be **position** and p be **momentum**

1. $p \sim \exp(K(p))$
2. $(q_i, p_i) \rightarrow (q_i^*, p_i^*)$ using Hamiltonian dynamics

$$\frac{dq_i}{dt} = \frac{\partial H}{\partial p_i} \quad \frac{dp_i}{dt} = -\frac{\partial H}{\partial q_i}$$

Hamiltonian motion's reversibility property $\implies \tilde{q}(p, q|p^*, q^*) = \tilde{q}(p^*, q^*|p, q)$

- ▶ conservation of Energy:

$$\begin{aligned} U(q) + K(p) &\approx U(q^*) + K(p^*) \\ \implies (p(q, p) \propto \exp(U(q) + K(p))) &\approx (p(q^*, p^*) \propto \exp(U(q^*) + K(p^*))) \\ \implies \min \left(1, \frac{\exp(U(q^*) + K(p^*))}{\exp(U(q) + K(p))} \frac{\tilde{q}(p, q|p^*, q^*)}{\tilde{q}(p^*, q^*|p, q)} \right) &\approx 1 \end{aligned}$$

- ▶ small rejection rate due to numerical errors in leap-frog

- Potts Model:


$$M(\Pi) \propto \exp \left(\sum_{i < j} \beta_{ij} \mathbf{1}_{z_i = z_j} \right)$$

- introduce auxiliary variables:

$$\Pr(r_{ij} = 0 | \Pi) = \exp(-\beta_{ij} \mathbf{1}_{z_i = z_j}) = q_{ij}$$

$$r_{ij} \sim \text{Bernoulli}(1 - \exp(-\beta_{ij} \mathbf{1}_{z_i = z_j}))$$

$$\Pr(\mathbf{r} | \Pi) = \prod_{1 \leq i < j \leq n} P(r_{ij} | \Pi)$$



- Conventional Swendsen-Wang and adding data \mathbf{y} : (see demo)

$$P(\Pi | \mathbf{r}, \mathbf{y}) \propto p(\mathbf{y} | \Pi) P(\mathbf{r} | \Pi) = \prod_{j=1}^k p(\mathbf{y}_{A_j}) \prod_{1 \leq i < j \leq n} \left[1 - \exp(-\beta_{ij} \mathbf{1}_{z_i = z_j}) \right]^{r_{ij}} \left[\exp(-\beta_{ij} \mathbf{1}_{z_i = z_j}) \right]^{1-r_{ij}}$$

- **key:** bond variables r_{ij} induce groups of sites which have same cluster label, a single $0^1 = 0$ makes the whole joint density become zero

► our past work

combine S-W with Bayesian Non-Parametrics:

Xu, R.Y.D. Caron, F., & Doucet, A (2016), Bayesian nonparametric image segmentation using a generalised Swendsen-Wang algorithm, arXiv:1602.03048

$$g(m_{-\ell,1}, \dots, m_{-\ell,j} + |C_\ell| \dots, m_{-\ell,k_\ell}) \frac{p(\mathbf{y}_{C_\ell \cup A_{-\ell,j}})}{p(\mathbf{y}_{A_{-\ell,j}})} \prod_{\{(i,j) | i \in C_\ell, r_{ij}=0\}} \exp\left(\beta_{ij}(1 - \delta_{ij})\mathbf{1}_{z_i=z_j}\right)$$

► current exploratory work:

a neural network module to learn better mixing rate

2.1.1 **Explain** reason why infinite-width Neural Network is a Gaussian Process

- \mathcal{GP} is a (potentially infinite) collection of RVs, s.t., joint distribution of every finite subset of RVs is multivariate Gaussian:

$$f \sim \mathcal{GP}(\mu(x), \mathcal{K}(x, x')) \quad \text{for any arbitrary } x, x'$$

- **prior** defined over $p(f|\mathcal{X})$, instead of $p(x)$ over $\mathcal{X} \equiv \{x_1, \dots, x_k\}$

$$p(f|\mathcal{X}) \equiv p\left(\begin{bmatrix} f(x_1) \\ \vdots \\ f(x_k) \end{bmatrix}\right) = \mathcal{N}(0, K) = \mathcal{N}\left(0, \begin{bmatrix} k(x_1, x_1) & \dots & k(x_1, x_k) \\ \vdots & \ddots & \vdots \\ k(x_k, x_1) & \dots & k(x_k, x_k) \end{bmatrix}\right)$$

marginal and conditional marginal under **noisy output**

- in a regression setting:

$$y_i = f(x_i) + \epsilon_i \quad \epsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma_\epsilon^2)$$

- joint $[\mathcal{Y}, y^*]^\top$, after integrate out f :

$$\begin{aligned} p\left(\begin{bmatrix} \mathcal{Y} \\ y^* \end{bmatrix} \middle| \begin{bmatrix} \mathcal{X} \\ x^{*\top} \end{bmatrix}, \sigma_\epsilon^2\right) &= \int p\left(\begin{bmatrix} \mathcal{Y} \\ y^* \end{bmatrix} \middle| \begin{bmatrix} \mathcal{X} \\ x^{*\top} \end{bmatrix}, f\right) p(f|\mathcal{X}, x^*) df \\ &= \int \mathcal{N}\left(\begin{bmatrix} \mathcal{Y} \\ y^* \end{bmatrix} \middle| \begin{bmatrix} f(\mathcal{X}) \\ f(x^{*\top}) \end{bmatrix}, \sigma_\epsilon^2 I\right) p(f|\mathcal{X}, x^*) df \\ &= \mathcal{N}\left(0, \begin{bmatrix} \underbrace{K(\mathcal{X}, \mathcal{X}) + \sigma_\epsilon^2 I}_{\Sigma_{1,1}} & \underbrace{K(\mathcal{X}, x^*)}_{\Sigma_{1,2}} \\ \underbrace{K(x^*, \mathcal{X})}_{\Sigma_{2,1}} & \underbrace{K(x^*, x^*) + \sigma_\epsilon^2}_{\Sigma_{2,2}} \end{bmatrix}\right) \end{aligned}$$

- predictive distribution** of $y^*|\mathcal{Y}$ using conditional formula of multivariate Gaussian:

$$\begin{aligned} p(y^*|\mathcal{Y}, \mathcal{X}, x^*) &= \mathcal{N}\left(\underbrace{\mathbf{0}}_{\mu_2} + \underbrace{K(x^*, \mathcal{X})}_{\Sigma_{2,1}} \underbrace{\left(K(\mathcal{X}, \mathcal{X}) + \sigma_\epsilon^2 I\right)^{-1}}_{\Sigma_{1,1}^{-1}} (\mathcal{Y} - \underbrace{\mathbf{0}}_{\mu_1}), \right. \\ &\quad \left. \underbrace{k(x^*, x^*) + \sigma_\epsilon^2}_{\Sigma_{2,2}} - \underbrace{K(x^*, \mathcal{X})}_{\Sigma_{2,1}} \underbrace{\left(K(\mathcal{X}, \mathcal{X}) + \sigma_\epsilon^2 I\right)^{-1}}_{\Sigma_{1,1}^{-1}} \underbrace{K(\mathcal{X}, x^*)}_{\Sigma_{1,2}}\right) \end{aligned}$$

marginal and conditional marginal under **noiseless output**

- **posterior** of f given \mathcal{Y} in regression:

$$p\left(\begin{bmatrix} \mathcal{Y} \\ f \end{bmatrix} \middle| \begin{bmatrix} \mathcal{X} \\ \mathbf{x}^\top \end{bmatrix}\right) = p\left(\begin{bmatrix} f(\mathcal{X}) \\ f(\mathbf{x}) \end{bmatrix}\right) = \mathcal{N}\left(0, \begin{bmatrix} K(\mathcal{X}, \mathcal{X}) + \sigma_\epsilon^2 \mathbf{I} & K(\mathcal{X}, \mathbf{x}) \\ K(\mathbf{x}, \mathcal{X}) & K(\mathbf{x}, \mathbf{x}) \end{bmatrix}\right)$$

for arbitrary variable \mathbf{x}

conditional marginal of $y^* | \mathcal{Y}$ using conditional formula of multivariate Gaussian:

$$p(f | \mathcal{X}, \mathcal{Y}) = \mathcal{GP}\left(K(\mathbf{x}, \mathcal{X})(K(\mathcal{X}, \mathcal{X}) + \sigma_\epsilon^2 \mathbf{I})^{-1} \mathcal{Y},\right. \\ \left. k(\mathbf{x}, \mathbf{x}') - K(\mathbf{x}, \mathcal{X}) \left(K(\mathcal{X}, \mathcal{X}) + \sigma_\epsilon^2 \mathbf{I}\right)^{-1} K(\mathcal{X}, \mathbf{x}')\right)$$

- **deterministic function** $y_i = f(x_i)$ is used, e.g., neural network's read-out layer $f(x_i)$, data y_i $p([\mathcal{Y}, y^*]^\top)$ no longer need to integrate f :

$$p\left(\begin{bmatrix} \mathcal{Y} \\ y^* \end{bmatrix} \middle| \begin{bmatrix} \mathcal{X} \\ \mathbf{x}^{*\top} \end{bmatrix}\right) = p\left(\begin{bmatrix} f(\mathcal{X}) \\ f(\mathbf{x}^*) \end{bmatrix}\right) = \mathcal{N}\left(0, \begin{bmatrix} K(\mathcal{X}, \mathcal{X}) & K(\mathcal{X}, \mathbf{x}^*) \\ K(\mathbf{x}^*, \mathcal{X}) & K(\mathbf{x}^*, \mathbf{x}^*) \end{bmatrix}\right)$$

predictive distribution $y^* | \mathcal{Y}$ using conditional formula of multivariate Gaussian:

$$p(y^* | \mathcal{Y}, \mathcal{X}, \mathbf{x}^*) = \mathcal{N}\left(K(\mathbf{x}^*, \mathcal{X})K(\mathcal{X}, \mathcal{X})^{-1} \mathcal{Y},\right. \\ \left. k(\mathbf{x}^*, \mathbf{x}^*) - K(\mathbf{x}^*, \mathcal{X})K(\mathcal{X}, \mathcal{X})^{-1} K(\mathcal{X}, \mathbf{x}^*)\right)$$

- ▶ proven by
"R. M. Neal. *Bayesian Learning for Neural Networks*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 1996"
and "J. H. Lee, Y. Bahri, R. Novak, S. S. Schoenholz, J. Pennington, and J. Sohl-Dickstein. *Deep neural networks as gaussian processes*. *ICLR*, 2018"
- ▶ at initialization:

$$z_k^l(x) = b_k^l + \sum_{j=1}^{N_l} w_{k,j}^l \times \phi\left(z_j^{l-1}(x)\right) \quad w_{k,j}^l \sim \mathcal{N}\left(0, \frac{1}{\sqrt{N_l}}\right) \quad b_k^l \sim \mathcal{N}(0, \sigma_b)$$

- ▶ interesting finding:

$$N_l \rightarrow \infty \implies z^l(\mathcal{X}) \sim \mathcal{GP}(0, \mathcal{K}^l(\mathcal{X}, \mathcal{X}))$$

with the recursive fomula: **NNGP**

$$K^l(x^{(p)}, x^{(q)}) = \sigma_b^2 + \sigma_w^2 \mathbb{E}_{(z_j^{l-1}(x^{(p)}), z_j^{l-1}(x^{(q)}))} \sim \mathcal{N}(0, K^{l-1}(x^{(p)}, x^{(q)})) \left[\phi(z_j^{l-1}(x^{(p)})) \phi(z_j^{l-1}(x^{(q)})) \right]$$

- ▶ **NNGP** was precursor towards **Neural Tangent Kernel**
my tutorial “Infinite Width: Neural Networks as Gaussian Process and Neural Tangent Kernel (NTK)”, https://github.com/roboticcam/machine-learning-notes/blob/master/files/gp_nn.pdf

- ▶ our work:

Huang, W., **Xu, R. Y. D.**, Du, W., Zeng Y., and Zhao Y., (2020) Mean field theory for deep dropout networks: digging up gradient backpropagation deeply, the 24th European Conference on Artificial Intelligence (ECAI 2020)

Huang, W., Du, W., **Xu, R. Y. D.**, (2020) On the Neural Tangent Kernel of Deep Networks with Orthogonal Initialization, arXiv preprint arXiv:2004.05867

2.2.1 Bayesian assisted **Generative Adversarial Network**

- ▶ GAN objective:

$$\begin{aligned}\min_G \max_D \left(\mathcal{L}(D, G) \equiv \mathbb{E}_{\mathbf{x} \sim p_r(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \right) \\ = \mathbb{E}_{\mathbf{x} \sim p_r(\mathbf{x})} [\log D(\mathbf{x})] + \underbrace{\mathbb{E}_{x \sim p_g(x)} [\log(1 - D(\mathbf{x}))]}_{\text{alternative expression}}\end{aligned}$$

- ▶ **my tutorial** on “Mathematics on Generative Adversarial Networks”:
<https://github.com/roboticcam/machine-learning-notes/blob/master/files/GAN.pdf>

- ▶ **Generator** $p(\theta_g | \mathbf{z}, \theta_d) \propto \left(\prod_{i=1}^{n_g} D_{\theta_d} \left(G_{\theta_g}(\mathbf{z}^{(i)}) \right) \right) p(\theta_g | \alpha_g)$
- ▶ **Discriminator** $p(\theta_d | \mathbf{z}, \mathbf{X}, \theta_g) \propto \prod_{i=1}^{n_d} D_{\theta_d}(\mathbf{x}^{(i)}) \times \prod_{i=1}^{n_g} \left(1 - D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)})) \right) \times p(\theta_d | \alpha_d)$
- ▶ **marginalization** $p(\theta_g | \theta_d)$

$$\begin{aligned}
 p(\theta_g | \theta_d) &= \int p(\theta_g, \mathbf{z} | \theta_d) d\mathbf{z} = \int p(\theta_g | \mathbf{z}, \theta_d) \underbrace{p(\mathbf{z} | \theta_d)}_{\text{independent of } \theta_d} d\mathbf{z} \\
 &= \int p(\theta_g | \mathbf{z}, \theta_d) p(\mathbf{z}) d\mathbf{z} \approx \frac{1}{N} \sum_{i=1}^N p(\theta_g | \mathbf{z}^{(i)}, \theta_d) \quad \mathbf{z}^{(i)} \sim p(\mathbf{z})
 \end{aligned}$$

- ▶ **marginalization** $p(\theta_d | \theta_g)$

$$\begin{aligned}
 p(\theta_d | \theta_g) &\equiv p(\theta_d | \mathbf{X}, \theta_g) = \int_{\mathbf{z}} p(\theta_d, \mathbf{z} | \mathbf{X}, \theta_g) d\mathbf{z} = \int p(\theta_d | \mathbf{z}, \mathbf{X}, \theta_g) \underbrace{p(\mathbf{z} | \mathbf{X}, \theta_g)}_{\text{independent of } \theta_d} d\mathbf{z} \\
 &= \int_{\mathbf{z}} p(\theta_d | \mathbf{z}, \mathbf{X}, \theta_g) p(\mathbf{z}) d\mathbf{z} \approx \frac{1}{N} \sum_{i=1}^N p(\theta_d | \mathbf{z}^{(i)}, \mathbf{X}, \theta_g) \quad \mathbf{z}^{(i)} \sim p(\mathbf{z})
 \end{aligned}$$

Saatchi, Y. and Wilson, A. G. *Bayesian gan*. In *Advances in Neural Information Processing Systems*, pp. 3625–3634, 2017

- ▶ look at GAN again:

$$\min_G \max_D \left(\mathcal{L}(D, G) \equiv \mathbb{E}_{\mathbf{x} \sim p_r(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_z \sim \underbrace{p_z(z)}_{\text{simple distribution}} [\log(1 - D(G(z)))] \right)$$

- ▶ researchers tend to use simple $p_z(z) \sim \mathcal{N}(0, \mathbf{I})$, but very complex G
- ▶ **question** is can we instead use more complex $p(z)$ to:
 - ▶ recover underlying data distribution $p_z(z)$
 - ▶ relief G , get $p_z(z)$ to do more work!

Our current and on-going work on sophisticated $p_z(z)$ in GAN

1. **work 1** replaced $p(z)$ in terms of Mixture densities:

Huang, W., **Xu, R Y D**, Jiang, S., Liang, X.; Oppermann, I., *GAN-Based Gaussian Mixture Model Responsibility Learning*, accepted to International Conference on Pattern Recognition 2021

2. **work 2** recent work replacing $p(z)$ with **Dirichlet Process** (under review)


Huang, W., **Xu, R Y D**, Jiang, S., Liang, X.; Oppermann, I., *Dirichlet Process Mixture Model Learning via GAN*

3. **work 3** exploratory work
use Hamiltonian Monte Carlo to sample complex $p_z(z)$

2.2.1 Determinantal Point Process (DPP) assisted Deep Learning

What is DPP?

- ▶ so, if it's a probability model, what is its parameter?
- ▶ we can either use a marginal kernel K , or to use an L-ensemble L
- ▶ let's look at marginal kernel first:



How do we define a DPP?

- ▶ Start with a **marginal** distribution of subset A

$$\Pr(A \subseteq \mathbf{Y}) = \det(K_A)$$

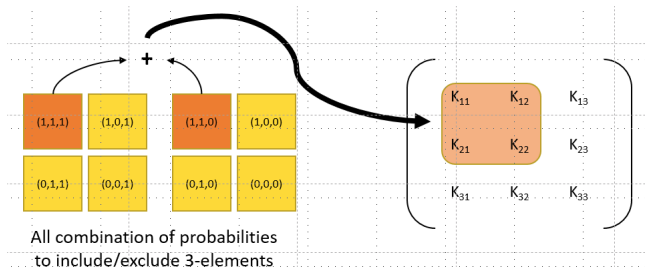
- ▶ An example: given $\mathbf{Y} = \{1, 2, 3, 4, 5\}$, $A = \{1, 2, 3\}$

$$\begin{aligned}\Pr(A \subseteq \mathbf{Y}) &\equiv \Pr(A \subseteq \mathbf{Y} \subseteq \mathbf{Y}) \equiv \Pr(\mathbf{Y} \in \{\{1, 2, 3\}, \{1, 2, 3, 4\}, \{1, 2, 3, 5\}, \{1, 2, 3, 4, 5\}\}) \\ &= \det(K_A)\end{aligned}$$

$$\begin{aligned}\Pr(A \subseteq \mathbf{Y}) &\equiv \Pr(A \subseteq \mathbf{Y} \subseteq \mathbf{Y}) \equiv \Pr(y_1 = 1, y_2 = 1, y_3 = 1) \\ &= \sum_{t_4=0}^1 \sum_{t_5=0}^1 \Pr(y_1 = 1, y_2 = 1, y_3 = 1, y_4 = t_4, y_5 = t_5) \\ &= \det(K_A)\end{aligned}$$

- ▶ **my tutorial** on “Determinantal Point Process”:
<https://github.com/roboticcam/machine-learning-notes/blob/master/files/dpp.pdf>

let's use a diagram!





- ▶ sum of probabilities to include 1st and 2nd elements
- ▶ This is defined by $\det(K_{\{1,2\}})$
- ▶ our past work on DPP:
 - ▶ Qiao, M., **Xu, R.Y.D.**, Bian, W. & Tao, D. (2016), 'Fast sampling for time-varying Determinantal Point Processes', vol. 11, no. 1. *ACM Transactions on Knowledge Discovery from Data*
 - ▶ Qiao, M., Bian, W., **Xu, R.Y.D.** & Tao, D. (2015), 'Diversified Hidden Markov Models for Sequential Labeling', vol. 27, no. 11, pp. 2947-2960. *IEEE Transactions on Knowledge and Data Engineering*



In Deep Learning: our work in Mini-batch selection via DPP

Huang, W., Xu, R.Y.D., I. Oppermann (2019), *Efficient Diversified Mini-Batch Selection using Variable High-layer Features*, *Asian Conference on Machine Learning, (ACML 2019)*

- Uniform sampling: less-diverse samples in one batch (lead to high variance in learning)



- ideally, each sample is a **good representation** of the entire data-set, higher diversities between selected samples (lead to low variance in learning)



Other example of Deep Learning assisted by DPP

- ▶ Z. Mariet and S. Sra. *Diversity networks: Neural network compression using Determinantal Point Processes*. In *International Conference on Learning Representations*, 2016.
- ▶ Zelda Mariet, Yaniv Ovadia, and Jasper Snoek. *Dppnet: Approximating determinantal point processes with deep networks*. *arXiv preprint arXiv:1901.02051*, 2019

2.2.3 probability distribution re-parameterization

Why Re-parameterization: (1) otherwise infeasible

- ▶ a **computation graph** of $L = f_3(f_2(f_1(\theta)))$, chain rule gives:

$$\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial f_3} \times \frac{\partial f_3}{\partial f_2} \times \frac{\partial f_2(\theta)}{\partial f_1(\theta)} \times \frac{\partial f_1}{\partial \theta} \quad f_2 = f_1(\theta)$$

- ▶ when f_2 is no longer deterministic:

$$\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial f_3} \times \frac{\partial f_3}{\partial f_2} \times \underbrace{\frac{\partial f_2(\theta)}{\partial f_1(\theta)}}_{\text{doesn't make sense!}} \times \frac{\partial f_1}{\partial \theta} \quad f_2 \sim f_1(\theta)$$

- ▶ solution, **re-parameterization trick**, s.t., random variable $\epsilon \sim q$ drawn from a distribution that is θ -free!

$$\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial f_3} \times \frac{\partial f_3}{\partial f_2} \times \underbrace{\frac{\partial f_2(\theta)}{\partial f_1(\theta)}}_{\text{OK again!}} \times \frac{\partial f_1}{\partial \theta} \quad f_2 = g(f_1(\theta), \epsilon) \quad \epsilon \sim q(\epsilon)$$

Simple Re-parameterization trick example

example: instead of:

$$z \sim \mathcal{N}(z; \mu(\theta), \sigma)$$

instead it can be re-parameterised into as a function of a standard Gaussian variable:

$$z = g(\epsilon, \theta) = \underbrace{\mu(\theta) + \epsilon\sigma}_{g(\epsilon, \theta)} \quad \epsilon \sim \underbrace{\mathcal{N}(0, 1)}_{p(\epsilon)}$$

Other re-parameterizations available?

- ▶ many available!

name	$p(z; \theta)$	$p(\epsilon)$	$g(\epsilon, \theta)$
Exponential	$\exp(-x); x > 0$	$\epsilon \sim [0; 1]$	$\ln(1/\epsilon)$
Cauchy	$\frac{1}{\pi(1+x^2)}$	$\epsilon \sim [0; 1]$	$\tan(\pi\epsilon)$
Laplace	$\mathcal{L}(0; 1) = \exp(- x)$	$\epsilon \sim [0; 1]$	$\ln(\frac{\epsilon_1}{\epsilon_2})$
Laplace	$\mathcal{L}(\mu; b)$	$\epsilon \sim [0; 1]$	$\mu - \text{bsgn}(\epsilon) \ln(1 - 2 \epsilon)$
Gaussian	$\mathcal{N}(0; 1)$	$\epsilon \sim [0; 1]$	$\sqrt{\ln(\frac{1}{\epsilon_1})} \cos(2\pi\epsilon_2)$
Gaussian	$\mathcal{N}(\mu; RR^\top)$	$\epsilon \sim \mathcal{N}(0; 1)$	$\mu + R\epsilon$
Rademacher	$\text{Rad}(\frac{1}{2})$	$\epsilon \sim \text{Bern}(\frac{1}{2})$	$2\epsilon - 1$
Log-Normal	$\ln \mathcal{N}(\mu; \sigma)$	$\epsilon \sim \mathcal{N}(\mu; \sigma^2)$	$\exp(\epsilon)$
Inv Gamma	$\mathcal{IG}(lk; \theta)$	$\epsilon \sim \mathcal{G}(k; \theta^{-1})$	$\frac{1}{\epsilon}$

- ▶ however, today we are interested only in **Softmax distribution** parameterizations!

Re-parameterization for variance reduction

- ▶ Re-parameterization is also used for variance reduction, discussing REBAR and RELAX algorithms:
- ▶ **my tutorial** on “Control Variate”:
https://github.com/roboticcam/machine-learning-notes/blob/master/files/variance_reduction.pdf
- ▶ will not discuss in this talk

- ▶ in deep learning, we need to perform:

$$k \sim \text{softmax}(\mu_1(\theta), \dots, \mu_L(\theta))$$

- ▶ μ_i defines **applications**:

- ▶ $\mu_i \equiv \mathbf{x}^\top \theta_i$ in classification
- ▶ $\mu_i \equiv \mathbf{u}_i^\top \mathbf{v}_c$ for word vectors

- ▶ but softmax function contains θ !, yes, in the chain rule path!

Re-parameterization using Gumbel-max trick

- ▶ instead of sample $k \sim \text{softmax}(\mu_1(\theta), \dots, \mu_K(\theta))$, we i.i.d. sample g instead
- ▶ we can perform **Gumbel-max** trick:

$$z_i \sim \text{ga}(z; 0, 1)$$

$$k = \arg \max_{i \in \{1, \dots, K\}} \{z_1 + \mu_1(\theta), \dots, z_K + \mu_K(\theta)\}$$

- ▶ well, there is two problems, firstly **why are the two are equivalent?**

Gumbel distribution definitions

- ▶ pdf of Gumbel with **unit scale** and location parameter μ :

$$p(Z = z|\mu, 1) \equiv \text{ga}(z; \mu) = \exp \left[- (z - \mu) - \exp(-(z - \mu)) \right]$$

- ▶ CDF of Gumbel:

$$\Pr(Z \leq z|\mu, 1) \equiv \mathcal{G}(z; \mu) = \exp \left[- \exp(-(z - \mu)) \right]$$

- ▶ it is obvious that:

$$\text{ga}(z|\mu, 1) = \exp(-z + \mu) \mathcal{G}(z|\mu)$$

which is a property you must know to work with Gumbels!

- ▶ some literature write location as $\log(\phi)$ instead of μ

$$\log(\phi) = \mu \quad \implies \quad \phi = \exp(\mu)$$

Gumbel-max trick and Softmax (1)

- ▶ given a set of Gumbel random variables $\{Z_i\}$, each having own location parameters $\{\mu_i\}$, probability of all other $Z_{i \neq k}$ are less than a particular value of z_k :

$$p(\max\{Z_{i \neq k}\} = z_k) = \prod_{i \neq k} \exp \left[-\exp\{-(z_k - \mu_i)\} \right]$$

- ▶ obviously, $Z_k \sim \text{gumbel}(Z_k = z_k; \mu_k)$:

$$\begin{aligned} & \Pr(k \text{ is largest} \mid \{\mu_i\}) \\ &= \int \exp\{-(z_k - \mu_k) - \exp\{-(z_k - \mu_k)\}\} \prod_{i \neq k} \exp\{-(z_k - \mu_i) - \exp\{-(z_k - \mu_i)\}\} dz_k \\ &= \int \exp \left[-z_k + \mu_k - \exp\{-(z_k - \mu_k)\} \right] \exp \left[-\sum_{i \neq k} \exp\{-(z_k - \mu_i)\} \right] dz_k \\ &= \int \exp \left[-z_k + \mu_k - \exp\{-(z_k - \mu_k)\} - \sum_{i \neq k} \exp\{-(z_k - \mu_i)\} \right] dz_k \\ &= \int \exp \left[-z_k + \mu_k - \sum_i \exp\{-(z_k - \mu_i)\} \right] dz_k \\ &= \int \exp \left[-z_k + \mu_k - \sum_i \exp\{-z_k + \mu_i\} \right] dz_k \\ &= \int \exp \left[-z_k + \mu_k - \exp\{-z_k\} \sum_i \exp\{\mu_i\} \right] dz_k \end{aligned}$$

Gumbel-max trick and Softmax (2)

► keep on going:

$$\begin{aligned}\Pr(k \text{ is largest} \mid \{\mu_i\}) &= \int \exp \left[-z_k + \mu_k - \exp\{-z_k\} \sum_i \exp\{\mu_i\} \right] dz_k \\&= \exp^{\mu_k} \int \exp \left[-z_k - \exp\{-z_k\} C \right] dz_k \\&= \exp^{\mu_k} \left[\frac{\exp(-C \exp(-z_k))}{C} \Big|_{z_k=-\infty}^{\infty} \right] \\&= \exp^{\mu_k} \left[\frac{1}{C} - 0 \right] \\&= \frac{\exp^{\mu_k}}{\sum_i \exp\{\mu_i\}}\end{aligned}$$

Gumbel-max trick summary

- ▶ moral of the story is, instead of sample from **softmax**

$$k \sim \left\{ \frac{\exp(\mu_1)}{\sum_i \exp(\mu_i)}, \dots, \frac{\exp(\mu_K)}{\sum_i \exp(\mu_i)} \right\}$$

- ▶ one can instead perform:

$$k = \arg \max_{i \in \{1, \dots, K\}} \{z_1, \dots, z_K\}$$

$$\text{where } z_i \sim \text{ga}(z; \mu_i) \equiv \exp \left[- (z - \mu_i) - \exp\{-(z - \mu_i)\} \right]$$

problem as μ still in Gumbel PDF, i.e., not “parameter-less” distribution

- ▶ using Gumbel's property of location:

$$k = \arg \max_{i \in \{1, \dots, K\}} \{\mu_1 + z_1, \dots, \mu_K + z_K\}$$

$$\text{where } z_i \stackrel{\text{iid}}{\sim} \text{ga}(z; 0) \equiv \exp \left[- (z) - \exp\{-(z)\} \right]$$

how to sample a Gumbel?

- ▶ CDF of a Gumbel, easy to take inverse, just keep “log”:

$$\begin{aligned}u &= \exp^{-\exp^{-(z-\mu)/\beta}} \\ \implies \log(u) &= -\exp^{-(z-\mu)/\beta} \\ \implies \log(-\log(u)) &= -(z-\mu)/\beta \\ \implies -\beta \log(-\log(u)) &= z-\mu \\ \implies z &= \mathcal{G}^{-1}(u) \equiv \mu - \beta \log(-\log(u))\end{aligned}$$

- ▶ for standard Gumbel, i.e., $\mu = 0, \beta = 1$:

$$z = \mathcal{G}^{-1}(u) \equiv -\log(-\log(u))$$

- ▶ therefore, sampling strategy:

$$\begin{aligned}u_i &\sim \mathcal{U}(0, 1) \\ z_i &= -\log(-\log(u_i)) \\ k &= \arg \max_{i \in \{1, \dots, K\}} \{\mu_1 + z_1, \dots, \mu_L + z_K\}\end{aligned}$$

Second problem with Softmax re-parameterisation

- ▶ look at re-parameterization:

$$z_i \sim \text{ga}(z; 0, 1)$$

$$k = \arg \max_{i \in \{1, \dots, K\}} \{z_1 + \mu_1(\theta), \dots, z_K + \mu_K(\theta)\}$$

- ▶ the other remaining **problem**: sample k also has an $\arg \max$ operation, it's a discrete distribution!
- ▶ one can **relax** the softmax distribution, for example **softmax map**
- ▶ several solutions proposed, for example:
"Maddison, Mnih, and Teh (2017), The Concrete Distribution: a Continuous Relaxation of Discrete Random Variables"

- ▶ use **softmax map** instead to propagate the gradient!

$$f_{\tau}(x)_k = \frac{\exp(\mu_k/\tau)}{\sum_{k=1}^K \exp(\mu_k/\tau)} \quad \mu_k \equiv \mu_k(x_k)$$

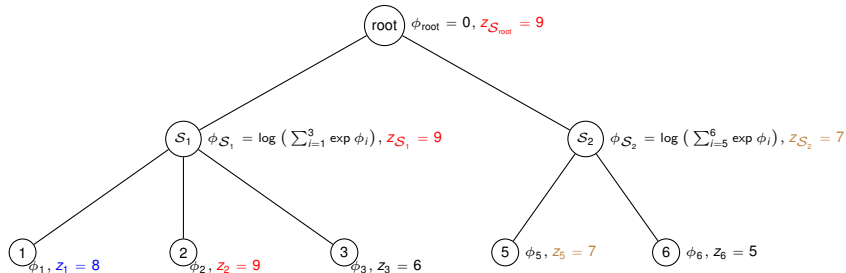
$$\text{as } \tau \rightarrow 0 \implies f_{\tau}(x) = \max \left(\left\{ \frac{\exp(\mu_k)}{\sum_{k=1}^K \exp(\mu_k)} \right\}_{k=1}^K \right)$$

- ▶ **questions** can you also think about the relationship between Gaussian Mixture Model and K-means?
- ▶ one can say $\tau = 1$ is softmax, and $\tau = 0$ is hard-max!
- ▶ then we can apply the same softmax map with added Gumbel variables:

$$z_i \sim \text{ga}(0)$$
$$f_{\tau}(\mu + \mathbf{z})_k = \frac{\exp(\mu_k + z_k)/\tau)}{\sum_{i=1}^K \exp(\mu_i + z_i)/\tau)}$$

- ▶ choose $\tau > 0$ for training, and $\tau = 0$ for inference

Sample Discrete Sequential use **Stochastic Beam Search**



$$\text{softmax}(\theta) \equiv \left(\underbrace{\frac{\exp(\mu_{\theta}(y_1|Y_{S_1}))}{\sum_{y'} \exp(\mu_{\theta}(y'_t|Y_{S_1}))}}_{\exp(\phi_1)}, \underbrace{\frac{\exp(\mu_{\theta}(y_2|Y_{S_1}))}{\sum_{y'} \exp(\mu_{\theta}(y'|Y_{S_1}))}}_{\exp(\phi_2)}, \underbrace{\frac{\exp(\mu_{\theta}(y_3|Y_{S_1}))}{\sum_{y'} \exp(\mu_{\theta}(y'|Y_{S_1}))}}_{\exp(\phi_3)} \right)$$

$$\phi_i = \log(p(y_i|Y_{S_i})) \equiv \mu_{\theta}(y_i|Y_{S_i}) - \log\left(\sum_{y'} \exp(\mu_{\theta}(y'|Y_{S_i}))\right)$$

Sample Discrete Sequential use **Stochastic Beam Search**

fact 1: $p_{\theta}(Y_S) = \underbrace{\exp(\phi_S)}_{\text{prob of node } S} = \sum_{i \in S} \underbrace{\exp(\phi_i)}_{\text{prob of child}} \implies \phi_S = \log p_{\theta}(Y_S) = \log \sum_{i \in S} \exp \phi_i$

fact 2: let $z_i \sim \text{ga}(z; \phi_i)$

max value: $\max_{i \in B} \{z_i\} \sim \text{ga}\left(\log \sum_{j \in B} \exp \phi_j\right)$

max index: $\arg \max_{i \in B} \{z_i\} \sim \text{Categorical}\left(\frac{\exp(\phi_i)}{\sum_{j \in B} \exp(\phi_j)}, i \in B\right)$

these two operations are independent, when combine the two, we have:

$$z_S = \max_{i \in S} \{z_i\} \sim \text{ga}\left(\underbrace{\log \sum_{i \in S} \exp \phi_i}_{\phi_S}\right) = \text{ga}(\phi_S)$$

fact 3: $p(\text{sample from softmax "without replacement" } m \text{ times}) =$
 $p(\text{choosing } m \text{ largest values from } \{z_i\})$

Top-down Stochastic Beams search

Wouter Kool, Herke van Hoof, and Max Welling. Stochastic beams and where to find them: The gumbel-top-k trick for sampling sequences without replacement. *arXiv preprint arXiv:1903.06059*, 2019

- ▶ Given a partial tree Y^S (we know its value of ϕ_S and Z_S):
 - ▶ For each $i \in \text{children}(S)$:

$$\phi_i \leftarrow \phi_S + \log p_\theta(Y^i | Y^S) : \quad \text{i.e., } \exp(\phi_i) = \exp(\phi_S) p_\theta(Y^i | Y^S)$$

S' extends partial tree length S by one token

$$z_i \sim \text{ga}(\phi_i)$$

- ▶ $Z = \max\{z_i\}$
- ▶ For each $i \in \text{children}(S)$:

$$\tilde{z}_i \leftarrow -\log(\exp(-Z_S) - \exp(-Z) + \exp(-z_i))$$

- ▶ $\text{BEAM} \leftarrow$ take top k of expansion according to $\{\tilde{z}_i\}$ then expand to add $(Y^i, \phi_i, \tilde{z}_i)$


understand $\tilde{z}_i \leftarrow -\log(\exp(-Z_S) - \exp(-Z) + \exp(-z_i))$ read original paper, or:

my tutorial on “Some New Research in Softmax”:

<https://github.com/roboticcam/machine-learning-notes/blob/master/files/softmax.pdf>

2.3.1 Long Short Term Memory to assist Kalman Filter's parameter training

- graphical model for state space model:



using **markov property** of probabilistic graphical model:

$$p(x_t | x_1, \dots, x_{t-1}, y_1, \dots, y_{t-1}) = p(x_t | x_{t-1})$$
$$p(y_t | x_1, \dots, x_{t-1}, x_t, y_1, \dots, y_{t-1}) = p(y_t | x_t)$$

- looks familiar to Recurrent Neural Networks?

What do we want to compute?

$$\text{Prediction : } p(x_t | \mathbf{y}_{1:t-1}) = \int_{x_{t-1}} p(x_t | x_{t-1}) p(x_{t-1} | \mathbf{y}_{1:t-1})$$

$$\text{Update : } p(x_t | \mathbf{y}_{1:t}) = \frac{p(y_t | x_t) p(x_t | \mathbf{y}_{1:t-1})}{\int_{s_t} p(y_t | s_t) p(ds_t | \mathbf{y}_{1:t-1})}$$

This is because:

$$\begin{aligned} p(x_t | \mathbf{y}_{1:t}) &\propto p(x_t, \mathbf{y}_{1:t}) \\ &\propto p(y_t | x_t) p(x_t | \mathbf{y}_{1:t-1}) \\ &= \frac{p(y_t | x_t) p(x_t | \mathbf{y}_{1:t-1})}{\int_{s_t} p(y_t | s_t) p(ds_t | \mathbf{y}_{1:t-1})} \end{aligned}$$

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B + w_t \quad w_t \sim \mathcal{N}(0, Q_t)$$

$$\Rightarrow \text{Transition probability:} \quad p(\mathbf{x}_t | \mathbf{x}_{t-1}) \sim \mathcal{N}(A\mathbf{x}_{t-1} + B, Q_t)$$

$$y_t = H\mathbf{x}_t + v_t \quad v_t \sim \mathcal{N}(0, R_t)$$

$$\Rightarrow \text{Measurement probability:} \quad p(y_t | \mathbf{x}_t) \sim \mathcal{N}(H\mathbf{x}_t, R_t)$$

- ▶ Kalman Filter can be used to in this Gaussian, Linear case.
- ▶ In general, there are many other Dyanmic models which are non-Gaussian, non-Linear. They can NOT be solved using Kalman Filter.

- ▶ Kalman filters require a motion model and measurement model to be specified at priory
- ▶ it's hard!
- ▶ can be crude approximation of reality
- ▶ this is where LSTM can help out!
- ▶ *H. Coskun, F. Achilles, R. DiPietro, N. Navab and F. Tombari, "Long short-term memory kalman filters: Recurrent neural estimators for pose regularization", Proc. IEEE Int. Conf. Comput. Vis., pp. 5525-5533, 2017*

- ▶ you see it everywhere, so I don't talk about it in detail:
- ▶ a compact form of representation:

$$\begin{bmatrix} i \\ f \\ o \\ \tilde{C} \end{bmatrix} = \begin{bmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{bmatrix} \left(\mathbf{W} \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix} + \mathbf{b} \right) \quad C_t = f_t \odot C_{t-1} + i \odot \tilde{C}_t$$

$$h_t = o_t \odot \tanh(C_t)$$


- ▶ in LSTMs, cell state C_t . The derivative of consecutive states is of the form:

$$\begin{aligned} C_t &= f_t(C_{t-1}) \odot C_{t-1} + i_t(C_{t-1}) \odot \tilde{C}_t(C_{t-1}) \\ &= f_t \odot C_{t-1} + i_t \odot \tanh(W_C[h_{t-1}, x_t] + b_C) \\ &= f_t(C_{t-1})C_{t-1} + \underbrace{i_t(h_{t-1}(C_{t-1}))\tanh(W_C[o_{t-1}(h_{t-1}(C_{t-1})), x_t] + b_C)}_{\xi(C_{t-1})} \end{aligned}$$

$$\frac{\partial C_t}{\partial C_{t-1}} = \underbrace{f_t}_{\text{gradient highway}} + \underbrace{\frac{\partial f_t}{\partial C_{t-1}} C_{t-1} + \frac{\partial \xi(C_{t-1})}{\partial C_{t-1}} C_{t-1}}_{\text{contains exponentially fast decay function}}$$

- ▶ of course, f_t may still close to zero
- ▶ **trick is to** initialize bias to be a large positive number, e.g.,
 $f_t = \sigma(W_f[h_{t-1}, x_t] + \text{large positive number})$ so to make f_t closer to 1 initially

LSTM Kalman Filters: Prediction




- the state-space model is changed into:

$$\begin{aligned} \mathbf{x}_t &= A\mathbf{x}_{t-1} + \mathbf{w}_t & \mathbf{w}_t &\sim \mathcal{N}(0, Q_t) & \rightarrow & \mathbf{x}_t = \text{lstm}_A(\mathbf{x}_{t-1}) + \mathbf{w}_t & \mathbf{w}_t &\sim \mathcal{N}(0, Q_t) \\ \mathbf{y}_t &= H\mathbf{x}_t + \mathbf{v}_t & \mathbf{v}_t &\sim \mathcal{N}(0, R_t) & \rightarrow & \text{unchanged} \end{aligned}$$

- prediction**

$$\begin{aligned} \bar{\mu}_t &= A\hat{\mu}_{t-1} & \rightarrow & \bar{\mu}_t = \text{lstm}_A(\hat{\mu}_{t-1}) \\ \bar{\Sigma}_t &= A\hat{\Sigma}_{t-1}A^T + Q_t & \rightarrow & \bar{\Sigma}_t = \mathcal{A}'\hat{\Sigma}_{t-1}\mathcal{A}'^T + \text{lstm}_Q(\bar{\mu}_t), \text{ where } \mathcal{A}' = \nabla_{\mathbf{x}_{t-1}} \text{lstm}_A(\hat{\mu}_{t-1}) \end{aligned}$$

LSTM Kalman Filters: Update



- the state-space model is changed into:

$$\begin{aligned} \mathbf{x}_t &= A\mathbf{x}_{t-1} + w_t & w_t &\sim \mathcal{N}(0, Q_t) & \rightarrow & \mathbf{x}_t = \text{lstm}_A(\mathbf{x}_{t-1}) + w_t & w_t &\sim \mathcal{N}(0, Q_t) \\ \mathbf{y}_t &= H\mathbf{x}_t + v_t & v_t &\sim \mathcal{N}(0, R_t) & \rightarrow & \text{unchanged} \end{aligned}$$

- update:

$$\begin{aligned} K_t &= \bar{\Sigma}_t H^T (H(\bar{\Sigma}_t)H^T + R_t)^{-1} & \rightarrow & K_t = \bar{\Sigma}_t H^T (H(\bar{\Sigma}_t)H^T + \text{lstm}_R(\mathbf{y}_t))^{-1} \\ \hat{\mu}_t &= \bar{\mu}_t + K(\mathbf{y}_t - H\bar{\mu}_t) & \rightarrow & \text{unchanged} \\ \hat{\Sigma}_t &= (I - KH)\bar{\Sigma}_t & \rightarrow & \text{unchanged} \end{aligned}$$

- ▶ Bayesian inference as a **research topic** will continue to play a key role in machine learning and many applicable applications.
- ▶ Bayesian inference as a **tool** will be able to enhance the functions of many machine learning models.

Thank you!